

## PRACTICAL 1

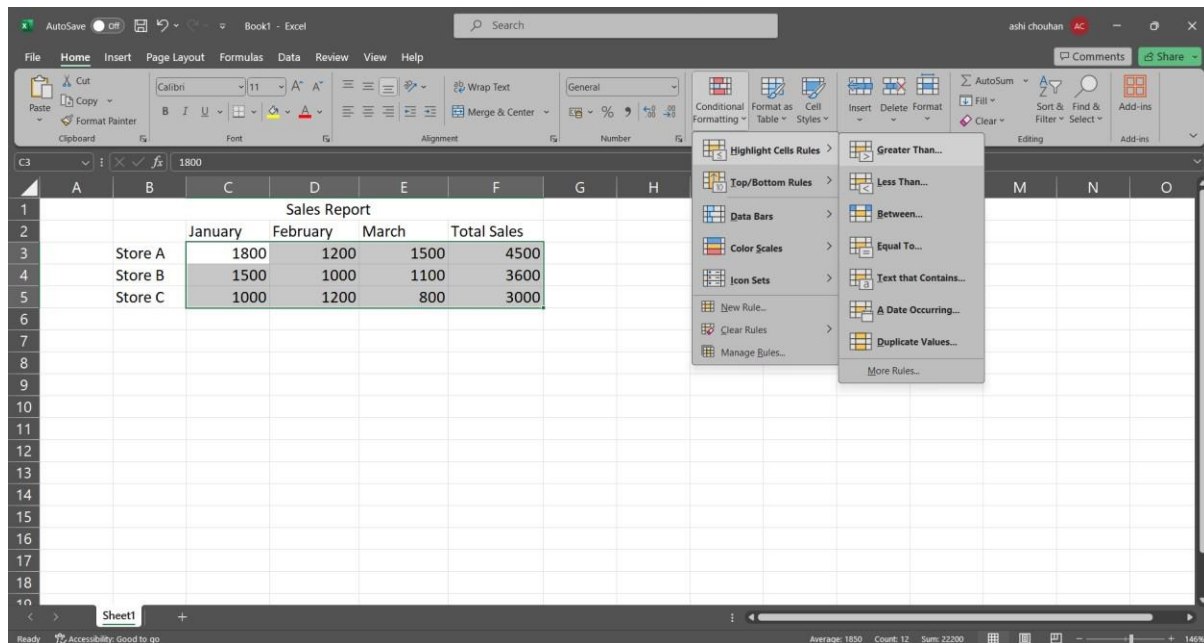
### Introduction to Excel

A. Perform conditional formatting on a dataset using various criteria.

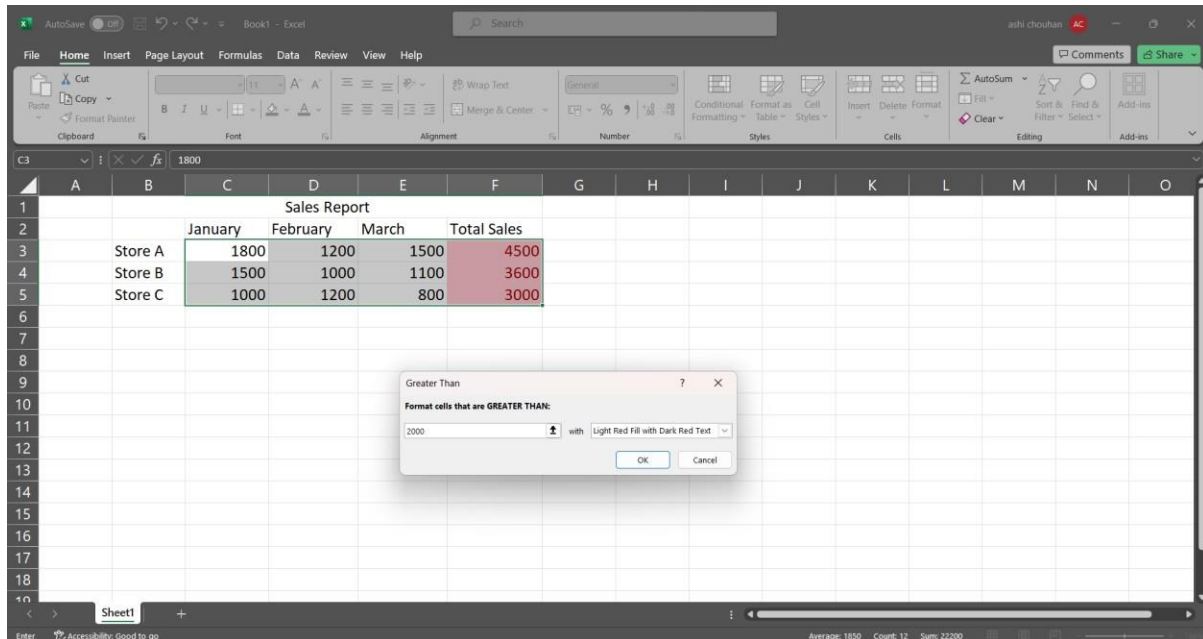
	A	B	C	D	E	F	G	H
1			Sales Report					
2			January	February	March	Total Sales		
3		Store A	1800	1200	1500	4500		
4		Store B	1500	1000	1100	3600		
5		Store C	1000	1200	800	3000		
6								
7								
8								
9								
10								
11								
12								
13								

### Steps

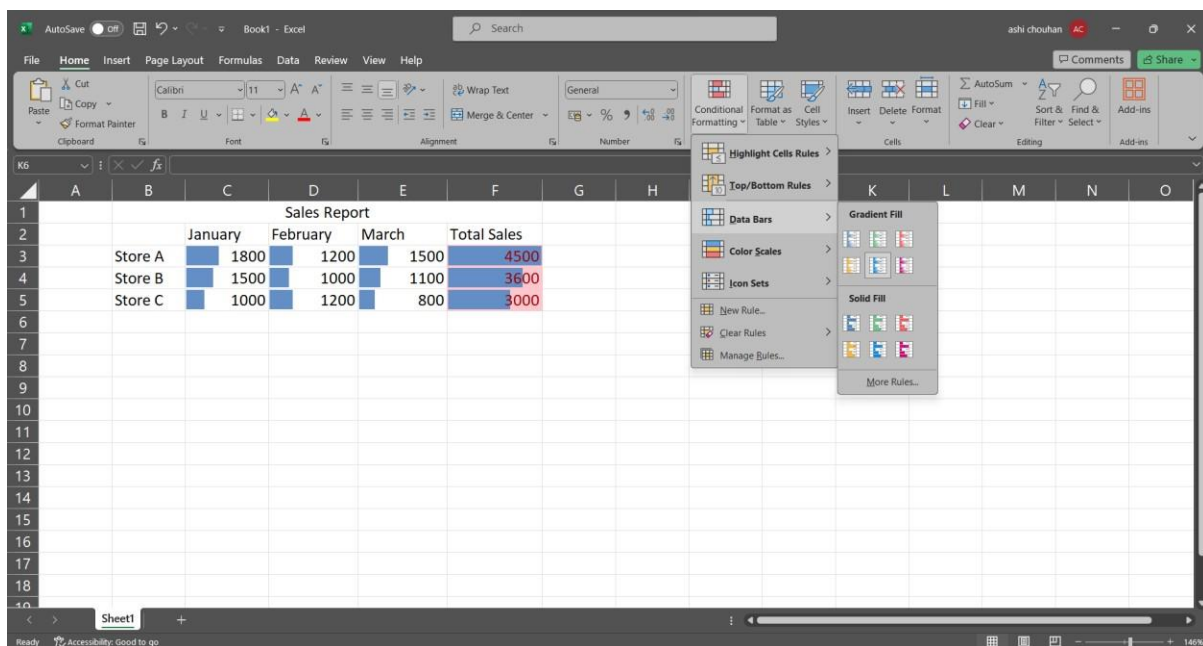
Step 1: Go to conditional formatting > Greater Than



Step 2: Enter the greater than filter value for example 2000.



Step 3: Go to Data Bars > Solid Fill in conditional formatting.

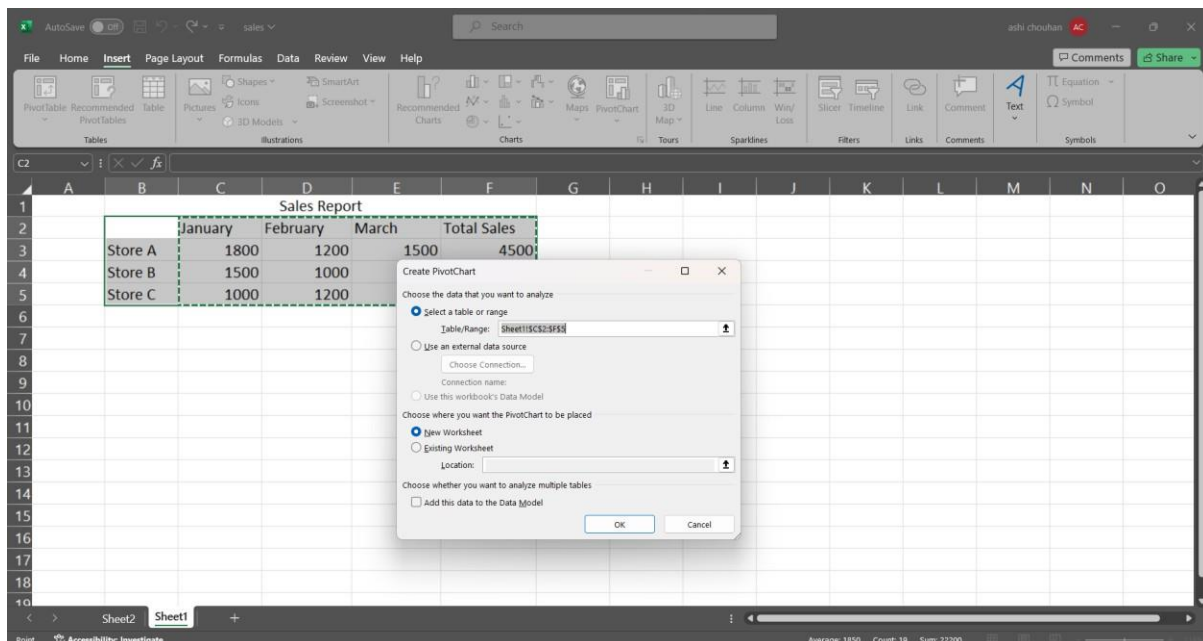


B. Create a pivot table to analyse and summarize data.

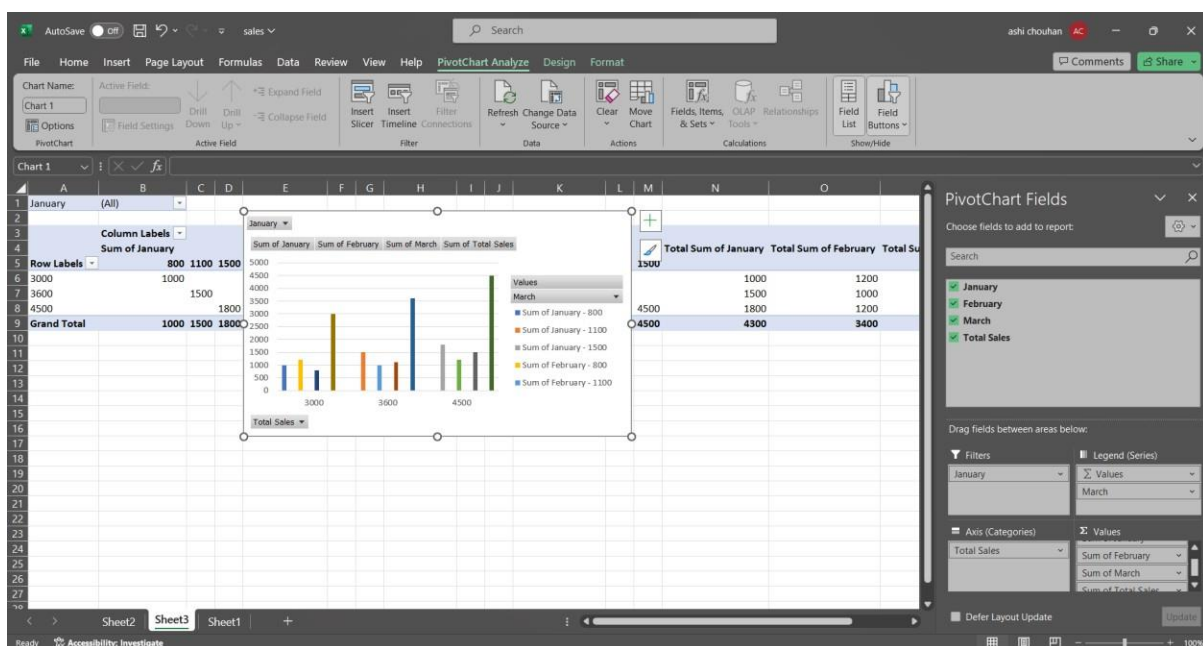
### Steps

Step 1: select the entire table and go to Insert tab PivotChart > Pivotchart

Step 2: Select “New worksheet” in the create pivot chart window.



Step 3: Select and drag attributes in the below boxes.

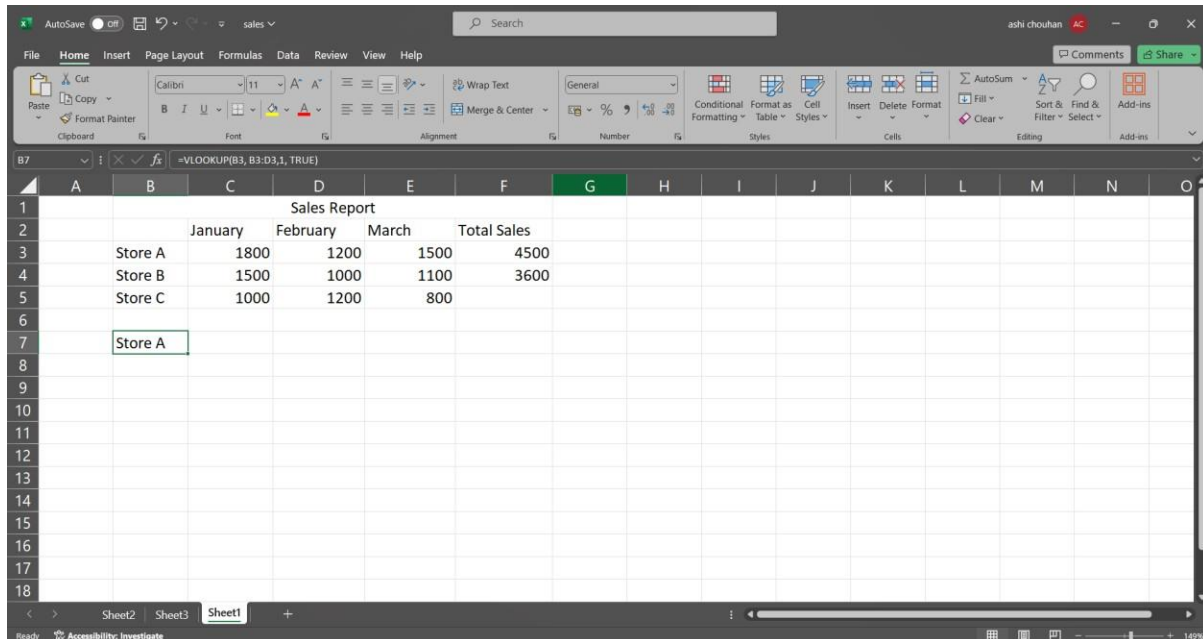


C. Use VLOOKUP function to retrieve information from a different worksheet or table.

Steps:

Step 1: click on an empty cell and type the following command.

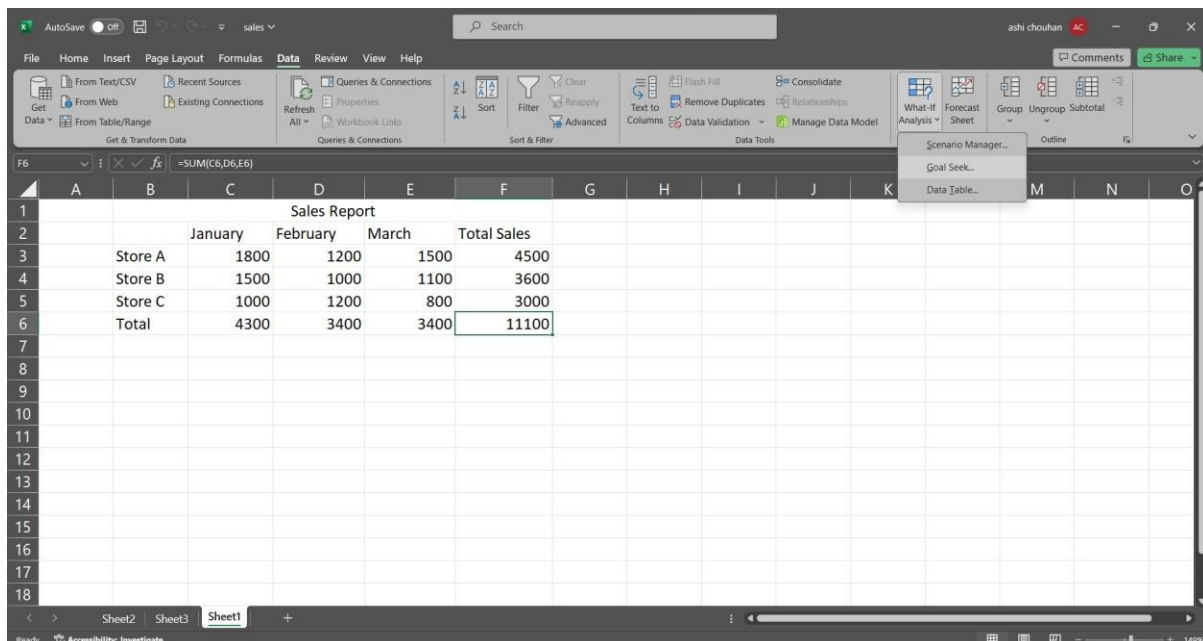
`=VLOOKUP(B3, B3:D3,1, TRUE)`



D. Perform what-if analysis using Goal Seek to determine input values for desired output.

Steps-

Step 1: In the Data tab go to the what if analysis>Goal seek.



Step 2: Fill the information in the window accordingly and click ok.

AutoSave sales

File Home Insert Page Layout Formulas Data Review View Help

Get Data From Text/CSV Recent Sources From Web Existing Connections From Table/Range

Get & Transform Data

Queries & Connections Refresh All Properties Workbook Links

Sort Filter Clear Reapply Advanced

Sort & Filter

Test to Columns Flash Fill Remove Duplicates Relationships What-If Analysis Forecast Sheet Group Ungroup Subtotal

Data Tools

Forecast

Outline

Comments Share

F6 =SUM(C6,D6,E6)

Sales Report					
	January	February	March	Total Sales	
Store A	1800	1200	1500	4500	
Store B	1500	1000	1100	3600	
Store C	1000	1200	800	3000	
Total	4300	3400	3400	11100	

Goal Seek

Set cell: F6

To value: 12000

By changing cell: E6

OK Cancel

AutoSave sales

File Home Insert Page Layout Formulas Data Review View Help

Get Data From Text/CSV Recent Sources From Web Existing Connections From Table/Range

Get & Transform Data

Queries & Connections Refresh All Properties Workbook Links

Sort Filter Clear Reapply Advanced

Sort & Filter

Test to Columns Flash Fill Remove Duplicates Relationships What-If Analysis Forecast Sheet Group Ungroup Subtotal

Data Tools

Forecast

Outline

Comments Share

F6 =SUM(C6,D6,E6)

Sales Report					
	January	February	March	Total Sales	
Store A	1800	1200	1500	4500	
Store B	1500	1000	1100	3600	
Store C	1000	1200	109700	111900	
Total	4300	3400	112300	120000	

Goal Seek Status

Goal Seeking with Cell F6 found a solution.

Target value: 120000

Current value: 120000

OK Cancel

AutoSave Off sales

File Home Insert Page Layout Formulas Data Review View Help

Comments Share

Get Data From Text/CSV Recent Sources Queries & Connections Sort Filter Clear Refresh All Properties Advanced Text to Columns Flash Fill Remove Duplicates Relationships What-If Analysis Forecast Sheet Group Ungroup Subtotal

Get & Transform Data Queries & Connections Sort & Filter Data Tools Manage Data Model Forecast Outline

F6 =SUM(C6:D6,E6)

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
1				Sales Report											
2			January	February	March	Total Sales									
3		Store A	1800	1200	1500	4500									
4		Store B	1500	1000	1100	3600									
5		Store C	1000	1200	109700	111900									
6		Total	4300	3400	112300	120000									
7															
8															
9															
10															
11															
12															
13															
14															
15															
16															
17															
18															

Ready Accessibility: Investigate

Sheet2 Sheet3 Sheet

## PRACTICAL 2

### Data Frames and Basic Data Pre-processing

A. Read data from CSV and JSON files into a data frame.

(1)

# Read data from a csv file

import pandas as pd

df = pd.read\_csv('Student\_Marks.csv')

print("Our dataset ")

print(df)

```
===== RESTART: D:\Notes\sem-6\data science\prac2
Our dataset
   number_courses  time_study  Marks
0                3      4.508  19.202
1                4      0.096   7.734
2                4      3.133  13.811
3                6      7.909  53.018
4                8      7.811  55.299
..             ...         ...    ...
95               6      3.561  19.128
96               3      0.301   5.609
97               4      7.163  41.444
98               7      0.309  12.027
99               3      6.335  32.357

[100 rows x 3 columns]
>>>
```

(2)

# Reading data from a JSON file

import pandas as pd

data = pd.read\_json('dataset.json')

print(data)

```
>>>
===== RESTART: D:/Notes/sem-6/data science/p:
   fruit  size  color
0  Apple  Large   Red
1  Banana Medium Yellow
2  Orange  Small  Orange
>>>
```

B. Perform basic data pre-processing tasks such as handling missing values and outliers.

Code:

(1)

# Replacing NA values using fillna()

import pandas as pd

```

df = pd.read_csv('titanic.csv')
print(df)
df.head(10)
print("Dataset after filling NA values with 0 : ")
df2=df.fillna(value=0)
print(df2)

```

```

===== RESTART: D:/Notes/sem-6/data science/prac2c.py =====
   PassengerId  Survived  Pclass    ... Cabin Embarked
0            892         0      3.0  ...   NaN         Q
1            893         0      3.0  ...   NaN         S
2            894         0      2.0  ...   NaN         Q
3            895         0      3.0  ...   NaN         S
4            896         0      3.0  ...   NaN         S
..          ...          ...      ...  ...   ...         ...
413          1305         0      3.0  ...   NaN         S
414          1306         0      1.0  ...  C105         C
415          1307         0      3.0  ...   NaN         S
416          1308         0      3.0  ...   NaN         S
417          1309         0      3.0  ...   NaN         C

[418 rows x 11 columns]
Dataset after filling NA values with 0 :
   PassengerId  Survived  Pclass    ... Cabin Embarked
0            892         0      3.0  ...     0         Q
1            893         0      3.0  ...     0         S
2            894         0      2.0  ...     0         Q
3            895         0      3.0  ...     0         S
4            896         0      3.0  ...     0         S
..          ...          ...      ...  ...   ...         ...
413          1305         0      3.0  ...     0         S
414          1306         0      1.0  ...  C105         C
415          1307         0      3.0  ...     0         S
416          1308         0      3.0  ...     0         S
417          1309         0      3.0  ...     0         C

[418 rows x 11 columns]

```

```

.>>>

```

(2)

```

# Dropping NA values using dropna()
import pandas as pd
df = pd.read_csv('titanic.csv')
print(df)
df.head(10)

print("Dataset after dropping NA values: ")
df.dropna(inplace = True)
print(df)

```



```

===== RESTART: D:/Notes/sem-6/data science/prac2c.py =====
   PassengerId  Survived  Pclass    ...    Cabin Embarked
0             892         3.0    ...    NaN        Q
1             893         3.0    ...    NaN        S
2             894         2.0    ...    NaN        Q
3             895         3.0    ...    NaN        S
4             896         NaN    ...    NaN        S
..          ...         ...    ...    ...        ...
413          1305         3.0    ...    NaN        S
414          1306         1.0    ...    C105        C
415          1307         3.0    ...    NaN        S
416          1308         3.0    ...    NaN        S
417          1309         3.0    ...    NaN        C

[418 rows x 11 columns]
Dataset after dropping NA values:
   PassengerId  Survived  Pclass    ...    Cabin Embarked
12             904         1.0    ...    B45        S
14             906         1.0    ...    E31        S
24             916         1.0    ...    B57 B59 B63 B66    C
26             918         1.0    ...    B36        C
28             920         1.0    ...    A21        S
..          ...         ...    ...    ...        ...
404          1296         1.0    ...    D40        C
405          1297         2.0    ...    D38        C
407          1299         1.0    ...    C80        C
411          1303         1.0    ...    C78        Q
414          1306         1.0    ...    C105        C

[87 rows x 11 columns]
>>>

```

C. Manipulate and transform data using functions like filtering, sorting, and grouping

Code:

```
import pandas as pd
```

```
# Load iris dataset
```

```
iris = pd.read_csv('Iris.csv')
```

```
# Filtering data based on a condition
```

```
setosa = iris[iris['Species'] == 'setosa']
```

```
print("Setosa samples:")
```

```
print(setosa.head())
```

```
# Sorting data
```

```
sorted_iris = iris.sort_values(by='SepalLengthCm', ascending=False)
```

```
print("\nSorted iris dataset:")
```

```
print(sorted_iris.head())
```

```
# Grouping data
```

```
grouped_species = iris.groupby('Species').mean()
```

```
print("\nMean measurements for each species:")
```

```
print(grouped_species)
```

```
===== RESTART: D:/Notes/sem-6/data science/prac2b.py =====
Setosa samples:
Empty DataFrame
Columns: [Id, SepalLengthCm, SepalWidthCm, PetalLengthCm, PetalWidthCm, Species]
Index: []

Sorted iris dataset:
   Id  SepalLengthCm  ...  PetalWidthCm  Species
131  132           7.9  ...           2.0  Iris-virginica
135  136           7.7  ...           2.3  Iris-virginica
122  123           7.7  ...           2.0  Iris-virginica
117  118           7.7  ...           2.2  Iris-virginica
118  119           7.7  ...           2.3  Iris-virginica

[5 rows x 6 columns]

Mean measurements for each species:
   Id  SepalLengthCm  ...  PetalLengthCm  PetalWidthCm
Species
Iris-setosa      25.5      5.006  ...      1.464      0.244
Iris-versicolor  75.5      5.936  ...      4.260      1.326
Iris-virginica  125.5      6.588  ...      5.552      2.026

[3 rows x 5 columns]
>>
```

## PRACTICAL 3

### Feature Scaling and Dummification

- A. Apply feature-scaling techniques like standardization and normalization to numerical features.

Code:

```
# Standardization and normalization
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler, StandardScaler
df = pd.read_csv('wine.csv', header=None, usecols=[0, 1, 2], skiprows=1)
df.columns = ['classlabel', 'Alcohol', 'Malic Acid']
print("Original DataFrame:")
print(df)
scaling=MinMaxScaler()
scaled_value=scaling.fit_transform(df[['Alcohol','Malic Acid']])
df[['Alcohol','Malic Acid']]=scaled_value
print("\n Dataframe after MinMax Scaling")
print(df)
scaling=StandardScaler()
scaled_standardvalue=scaling.fit_transform(df[['Alcohol','Malic Acid']])
df[['Alcohol','Malic Acid']]=scaled_standardvalue
print("\n Dataframe after Standard Scaling")
print(df)
```

```
= RESTART: D:/Notes/sem-6/data science/prac3b.py
```

```
Original DataFrame:
```

	classlabel	Alcohol	Malic Acid
0	1	14.23	1.71
1	1	13.20	1.78
2	1	13.16	2.36
3	1	14.37	1.95
4	1	13.24	2.59
..	...	...	...
173	3	13.71	5.65
174	3	13.40	3.91
175	3	13.27	4.28
176	3	13.17	2.59
177	3	14.13	4.10

```
[178 rows x 3 columns]
```

```
Dataframe after MinMax Scaling
```

	classlabel	Alcohol	Malic Acid
0	1	0.842105	0.191700
1	1	0.571053	0.205534
2	1	0.560526	0.320158
3	1	0.878947	0.239130
4	1	0.581579	0.365613
..	...	...	...
173	3	0.705263	0.970356
174	3	0.623684	0.626482
175	3	0.589474	0.699605
176	3	0.563158	0.365613
177	3	0.815789	0.664032

```
[178 rows x 3 columns]
```

```
Dataframe after Standard Scaling
```

173	3	0.305130	0.305010
177	3	0.815789	0.664032

```
[178 rows x 3 columns]
```

```
Dataframe after Standard Scaling
```

	classlabel	Alcohol	Malic Acid
0	1	1.518613	-0.562250
1	1	0.246290	-0.499413
2	1	0.196879	0.021231
3	1	1.691550	-0.346811
4	1	0.295700	0.227694
..	...	...	...
173	3	0.876275	2.974543
174	3	0.493343	1.412609
175	3	0.332758	1.744744
176	3	0.209232	0.227694
177	3	1.395086	1.583165

```
[178 rows x 3 columns]
```

```
>>>
```

B. Perform feature Dummification to convert categorical variables into numerical representations.

Code:

```
import pandas as pd
iris=pd.read_csv("Iris.csv")
print(iris)
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
iris['code']=le.fit_transform(iris.Species)
print(iris)
```

```
===== RESTART: D:/Notes/sem-6/data science/prac3a.py =====
   Id  SepalLengthCm  ...  PetalWidthCm  Species
0    1             5.1  ...           0.2  Iris-setosa
1    2             4.9  ...           0.2  Iris-setosa
2    3             4.7  ...           0.2  Iris-setosa
3    4             4.6  ...           0.2  Iris-setosa
4    5             5.0  ...           0.2  Iris-setosa
..  ...           ...  ...           ...  ...
145 146             6.7  ...           2.3  Iris-virginica
146 147             6.3  ...           1.9  Iris-virginica
147 148             6.5  ...           2.0  Iris-virginica
148 149             6.2  ...           2.3  Iris-virginica
149 150             5.9  ...           1.8  Iris-virginica

[150 rows x 6 columns]
   Id  SepalLengthCm  SepalWidthCm  ...  PetalWidthCm  Species  code
0    1             5.1           3.5  ...           0.2  Iris-setosa    0
1    2             4.9           3.0  ...           0.2  Iris-setosa    0
2    3             4.7           3.2  ...           0.2  Iris-setosa    0
3    4             4.6           3.1  ...           0.2  Iris-setosa    0
4    5             5.0           3.6  ...           0.2  Iris-setosa    0
..  ...           ...           ...  ...           ...  ...  ...
145 146             6.7           3.0  ...           2.3  Iris-virginica    2
146 147             6.3           2.5  ...           1.9  Iris-virginica    2
147 148             6.5           3.0  ...           2.0  Iris-virginica    2
148 149             6.2           3.4  ...           2.3  Iris-virginica    2
149 150             5.9           3.0  ...           1.8  Iris-virginica    2

[150 rows x 7 columns]
>>>
```

## Practical 4

### Hypothesis Testing

Conduct a hypothesis test using appropriate statistical tests (e.g., t-test, chi-square test)

```
# t-test
```

```
import numpy as np
```

```
from scipy import stats
```

```
import matplotlib.pyplot as plt
```

```
# Generate two samples for demonstration purposes
```

```
np.random.seed(42)
```

```
sample1 = np.random.normal(loc=10, scale=2, size=30)
```

```
sample2 = np.random.normal(loc=12, scale=2, size=30)
```

```
# Perform a two-sample t-test
```

```
t_statistic, p_value = stats.ttest_ind(sample1, sample2)
```

```
# Set the significance level
```

```
alpha = 0.05
```

```
print("Results of Two-Sample t-test:")
```

```
print(f'T-statistic: {t_statistic}')
```

```
print(f'P-value: {p_value}')
```

```
print(f'Degrees of Freedom: {len(sample1) + len(sample2) - 2}')
```

```
# Plot the distributions
```

```
plt.figure(figsize=(10, 6))
```

```
plt.hist(sample1, alpha=0.5, label='Sample 1', color='blue')
```

```

plt.hist(sample2, alpha=0.5, label='Sample 2', color='orange')
plt.axvline(np.mean(sample1), color='blue', linestyle='dashed', linewidth=2)
plt.axvline(np.mean(sample2), color='orange', linestyle='dashed', linewidth=2)
plt.title('Distributions of Sample 1 and Sample 2')
plt.xlabel('Values')
plt.ylabel('Frequency')
plt.legend()

# Highlight the critical region if null hypothesis is rejected
if p_value < alpha:
    critical_region = np.linspace(min(sample1.min(), sample2.min()), max(sample1.max(),
sample2.max()), 1000)
    plt.fill_between(critical_region, 0, 5, color='red', alpha=0.3, label='Critical Region')
    plt.text(11, 5, f'T-statistic: {t_statistic:.2f}', ha='center', va='center', color='black',
backgroundcolor='white')

# Show the plot
plt.show()

# Draw Conclusions
if p_value < alpha:
    if np.mean(sample1) > np.mean(sample2):
        print("Conclusion: There is significant evidence to reject the null hypothesis.")
        print("Interpretation: The mean of Sample 1 is significantly higher than that of Sample
2.")
    else:
        print("Conclusion: There is significant evidence to reject the null hypothesis.")
        print("Interpretation: The mean of Sample 2 is significantly higher than that of Sample
1.")
    else:


```

```
print("Conclusion: Fail to reject the null hypothesis.")
```

```
print("Interpretation: There is not enough evidence to claim a significant difference  
between the means.")
```

Output:

```
----- RESTART: Python 3.7.4 Shell -----  
Results of Two-Sample t-test:  
T-statistic: -4.512913234547555  
P-value: 3.176506547470154e-05  
Degrees of Freedom: 58
```



```
#chi-test
```

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib as plt
```

```
import seaborn as sb
```

```
import warnings
```



```
from scipy import stats
warnings.filterwarnings('ignore')
df=sb.load_dataset('mpg')
print(df)
print(df['horsepower'].describe())
print(df['model_year'].describe())
bins=[0,75,150,240]
df['horsepower_new']=pd.cut(df['horsepower'],bins=bins,labels=['l','m','h'])
c=df['horsepower_new']
print(c)
ybins=[69,72,74,84]
label=['t1','t2','t3']
df['modelyear_new']=pd.cut(df['model_year'],bins=ybins,labels=label)
newyear=df['modelyear_new']
print(newyear)
df_chi=pd.crosstab(df['horsepower_new'],df['modelyear_new'])
print(df_chi)
print(stats.chi2_contingency(df_chi))
Output:
```

```

----- RESIDENT: E:/all notes/DS/plac_4.1.py -----
      mpg  cylinders  ...  origin  name
0      18.0         8  ...    usa  chevrolet chevelle malibu
1      15.0         8  ...    usa      buick skylark 320
2      18.0         8  ...    usa    plymouth satellite
3      16.0         8  ...    usa      amc rebel sst
4      17.0         8  ...    usa      ford torino
..      ...         ...  ...    ...      ...
393    27.0         4  ...    usa      ford mustang gl
394    44.0         4  ... europe      vw pickup
395    32.0         4  ...    usa      dodge rampage
396    28.0         4  ...    usa      ford ranger
397    31.0         4  ...    usa      chevy s-10

[398 rows x 9 columns]
count      392.000000
mean       104.469388
std         38.491160
min         46.000000
25%         75.000000
50%         93.500000
75%        126.000000
max        230.000000

```

```
Name: horsepower, dtype: float64
count      398.000000
mean        76.010050
std          3.697627
min         70.000000
25%         73.000000
50%         76.000000
75%         79.000000
max         82.000000
```

```
Name: model_year, dtype: float64
0          m
1          h
2          m
3          m
4          m
..
393        m
394        l
395        m
396        m
397        m
```

```
Name: horsepower_new, Length: 398, dtype: category
Categories (3, object): ['l' < 'm' < 'h']
```

```
0      t1
1      t1
2      t1
3      t1
4      t1
..
393    t3
394    t3
395    t3
396    t3
397    t3
```

```
Name: modelyear_new, Length: 398, dtype: category
Categories (3, object): ['t1' < 't2' < 't3']
```

```
modelyear_new  t1  t2  t3
horsepower_new
```

```
l           9  14   76
m          49  41  158
h          26  11    8
```

```
(54.95485392447537, 3.320518009555984e-11, 4, array([[ 21.21428571,  16.66836735,  61.11734694]
,
[ 53.14285714,  41.75510204, 153.10204082],
[  9.64285714,   7.57653061,  27.78061224]]))
```

**Conclusion:** There is sufficient evidence to reject the null hypothesis, indicating that there is a significant association between 'horsepower\_new' and 'modelyear\_new' categories.

## Practical 5

### ANOVA (Analysis of Variance)

Perform one-way ANOVA to compare means across multiple groups.

Conduct post-hoc tests to identify significant differences between group means.

```
import pandas as pd
import scipy.stats as stats
from statsmodels.stats.multicomp import pairwise_tukeyhsd

group1 = [23, 25, 29, 34, 30]
group2 = [19, 20, 22, 24, 25]
group3 = [15, 18, 20, 21, 17]
group4 = [28, 24, 26, 30, 29]

all_data = group1 + group2 + group3 + group4
group_labels = ['Group1'] * len(group1) + ['Group2'] * len(group2) + ['Group3'] *
len(group3) + ['Group4'] * len(group4)

f_statistics, p_value = stats.f_oneway(group1, group2, group3, group4)
print("one-way ANOVA:")
print("F-statistics:", f_statistics)
print("p-value", p_value)

tukey_results = pairwise_tukeyhsd(all_data, group_labels)
print("\nTukey-Kramer post-hoc test:")
print(tukey_results)
```

Output:-

REPORT: F-TEST RESULTS/20/

one-way ANOVA:  
F-statistics: 12.139872842870115  
p-value 0.00021465200901629603

Tukey-Kramer post-hoc test:  
Multiple Comparison of Means - Tukey HSD, FWER=0.05

=====						
group1	group2	meandiff	p-adj	lower	upper	reject
-----						
Group1	Group2	-6.2	0.024	-11.6809	-0.7191	True
Group1	Group3	-10.0	0.0004	-15.4809	-4.5191	True
Group1	Group4	-0.8	0.9747	-6.2809	4.6809	False
Group2	Group3	-3.8	0.2348	-9.2809	1.6809	False
Group2	Group4	5.4	0.0542	-0.0809	10.8809	False
Group3	Group4	9.2	0.001	3.7191	14.6809	True
-----						

|

## **Practical 6**

### **Regression and its Types.**

```
import numpy as np
import pandas as pd
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

housing = fetch_california_housing()
housing_df = pd.DataFrame(housing.data, columns=housing.feature_names)
print(housing_df)

housing_df['PRICE'] = housing.target

X = housing_df[['AveRooms']]
y = housing_df['PRICE']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model = LinearRegression()

model.fit(X_train, y_train)

mse = mean_squared_error(y_test, model.predict(X_test))
r2 = r2_score(y_test, model.predict(X_test))
```

```
print("Mean Squared Error:", mse)
print("R-squared:", r2)
print("Intercept:", model.intercept_)
print("Coefficient:", model.coef_)
```

```
#####
```

```
#Multiple Liner Regression
```

```
X = housing_df.drop('PRICE',axis=1)
y = housing_df['PRICE']
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=42)
```

```
model = LinearRegression()
```

```
model.fit(X_train,y_train)
```

```
y_pred = model.predict(X_test)
```

```
mse = mean_squared_error(y_test,y_pred)
```

```
r2 = r2_score(y_test,y_pred)
```

```
print("Mean Squared Error:",mse)
print("R-squared:",r2)
print("Intercept:",model.intercept_)
print("Coefficient:",model.coef_)
```

Output:

```
----- RESIDENTIAL_E:\all notes\DS\prac_9_single.py -----
      MedInc  HouseAge  AveRooms  ...  AveOccup  Latitude  Longitude
0      8.3252      41.0  6.984127  ...  2.555556      37.88      -122.23
1      8.3014      21.0  6.238137  ...  2.109842      37.86      -122.22
2      7.2574      52.0  8.288136  ...  2.802260      37.85      -122.24
3      5.6431      52.0  5.817352  ...  2.547945      37.85      -122.25
4      3.8462      52.0  6.281853  ...  2.181467      37.85      -122.25
...      ...      ...      ...      ...      ...      ...
20635  1.5603      25.0  5.045455  ...  2.560606      39.48      -121.09
20636  2.5568      18.0  6.114035  ...  3.122807      39.49      -121.21
20637  1.7000      17.0  5.205543  ...  2.325635      39.43      -121.22
20638  1.8672      18.0  5.329513  ...  2.123209      39.43      -121.32
20639  2.3886      16.0  5.254717  ...  2.616981      39.37      -121.24

[20640 rows x 8 columns]
Mean Squared Error: 1.2923314440807299
R-squared: 0.013795337532284901
Intercept: 1.654762268596842
Coefficient: [0.07675559]
Mean Squared Error: 0.5558915986952441
R-squared: 0.575787706032451
Intercept: -37.02327770606414
Coefficient: [ 4.48674910e-01  9.72425752e-03 -1.23323343e-01  7.83144907e-01
 -2.02962058e-06 -3.52631849e-03 -4.19792487e-01 -4.33708065e-01]
```



## Practical 7

### Logistic Regression and Decision Tree

```
import numpy as np
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier

from sklearn.metrics import accuracy_score, precision_score, recall_score,
classification_report

# Load the Iris dataset and create a binary classification problem
iris = load_iris()

iris_df = pd.DataFrame(data=np.c_[iris['data'], iris['target']], columns=iris['feature_names'] +
['target'])

binary_df = iris_df[iris_df['target'] != 2]

X = binary_df.drop('target', axis=1)
y = binary_df['target']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train a logistic regression model and evaluate its performance
logistic_model = LogisticRegression()
logistic_model.fit(X_train, y_train)
y_pred_logistic = logistic_model.predict(X_test)

print("Logistic Regression Metrics")
print("Accuracy: ", accuracy_score(y_test, y_pred_logistic))
print("Precision:", precision_score(y_test, y_pred_logistic))
print("Recall: ", recall_score(y_test, y_pred_logistic))
```

```
print("\nClassification Report")
print(classification_report(y_test, y_pred_logistic))
# Train a decision tree model and evaluate its performance
decision_tree_model = DecisionTreeClassifier()
decision_tree_model.fit(X_train, y_train)
y_pred_tree = decision_tree_model.predict(X_test)
print("\nDecision Tree Metrics")
print("Accuracy: ", accuracy_score(y_test, y_pred_tree))
print("Precision:", precision_score(y_test, y_pred_tree))
print("Recall: ", recall_score(y_test, y_pred_tree))
print("\nClassification Report")
print(classification_report(y_test, y_pred_tree))
```

Output:-

```
Logistic Regression Metrics
Accuracy:  1.0
Precision: 1.0
Recall:    1.0
```

```
Classification Report
              precision    recall  f1-score   support

    0.0         1.00      1.00      1.00         12
    1.0         1.00      1.00      1.00          8

   accuracy                1.00         20
  macro avg              1.00      1.00      1.00         20
weighted avg              1.00      1.00      1.00         20
```

```
Decision Tree Metrics
Accuracy:  1.0
Precision: 1.0
Recall:    1.0
```

```
Classification Report
              precision    recall  f1-score   support

    0.0         1.00      1.00      1.00         12
    1.0         1.00      1.00      1.00          8

   accuracy                1.00         20
  macro avg              1.00      1.00      1.00         20
weighted avg              1.00      1.00      1.00         20
```

## Practical 8

### K-Means clustering

```
import pandas as pd

from sklearn.preprocessing import MinMaxScaler

from sklearn.cluster import KMeans

import matplotlib.pyplot as plt

data = pd.read_csv("C:\\Users\\Reape\\Downloads\\wholesale\\wholesale.csv")

data.head()

categorical_features = ['Channel', 'Region']

continuous_features = ['Fresh', 'Milk', 'Grocery', 'Frozen', 'Detergents_Paper', 'Delicassen']

data[continuous_features].describe()

for col in categorical_features:

    dummies = pd.get_dummies(data[col], prefix = col)

    data = pd.concat([data, dummies], axis = 1)

    data.drop(col, axis = 1, inplace = True)

data.head()

mms = MinMaxScaler()

mms.fit(data)

data_transformed = mms.transform(data)

sum_of_squared_distances = []

K = range(1, 15)

for k in K:

    km = KMeans(n_clusters=k)

    km = km.fit(data_transformed)
```

```
sum_of_squared_distances.append(km.inertia_)
```

```
plt.plot(K, sum_of_squared_distances, 'bx-')
```


```
plt.xlabel('k')
```

```
plt.ylabel('sum_of_squared_distances')
```

```
plt.title('elbow Mehtod for optimal k')
```

```
plt.show()
```

Output:



## Practical 9

### Principal Component Analysis (PCA)

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA

iris = load_iris()
iris_df = pd.DataFrame(data=np.c_[iris['data'], iris['target']], columns=iris['feature_names'] +
['target'])
X = iris_df.drop('target', axis=1)
y = iris_df['target']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

pca = PCA()
X_pca = pca.fit_transform(X_scaled)
explained_variance_ratio = pca.explained_variance_ratio_

plt.figure(figsize=(8, 6))
plt.plot(np.cumsum(explained_variance_ratio), marker='o', linestyle='--')
plt.title('Explained Variance Ratio')
plt.xlabel('Number of Principal Components')
plt.ylabel('Cumulative Explained Variance Ratio')
plt.grid(True)
```

```
plt.show()
```

```
cumulative_variance_ratio = np.cumsum(explained_variance_ratio)
```

```
n_components = np.argmax(cumulative_variance_ratio >= 0.95) + 1
```

```
print(f"Number of principal components to explain 95% variance: {n_components}")
```

```
pca = PCA(n_components=n_components)
```

```
X_reduced = pca.fit_transform(X_scaled)
```

```
plt.figure(figsize=(8, 6))
```

```
plt.scatter(X_reduced[:, 0], X_reduced[:, 1], c=y, cmap='viridis', s=50, alpha=0.5)
```

```
plt.title('Data in Reduced-dimensional Space')
```



```
plt.xlabel('Principal Component 1')
```

```
plt.ylabel('Principal Component 2')
```

```
plt.colorbar(label='Target')
```

```
plt.show()
```

Output:



```
----- RESTART: E.: all nodes/DS/pl
Number of principal components to explain 95% variance: 2
```



## Practical 10

### Data Visualization and Storytelling

```
import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns

import numpy as np

# Generate random data

np.random.seed(42) # Set a seed for reproducibility

# Create a DataFrame with random data

data = pd.DataFrame({

    'variable1': np.random.normal(0, 1, 1000),

    'variable2': np.random.normal(2, 2, 1000) + 0.5 * np.random.normal(0, 1, 1000),

    'variable3': np.random.normal(-1, 1.5, 1000),

    'category': pd.Series(np.random.choice(['A', 'B', 'C', 'D'], size=1000, p=[0.4, 0.3, 0.2, 0.1]),

dtype='category')

})

# Create a scatter plot to visualize the relationship between two variables

plt.figure(figsize=(10, 6))

plt.scatter(data['variable1'], data['variable2'], alpha=0.5)

plt.title('Relationship between Variable 1 and Variable 2', fontsize=16)

plt.xlabel('Variable 1', fontsize=14)

plt.ylabel('Variable 2', fontsize=14)

plt.show()
```

```
# Create a bar chart to visualize the distribution of a categorical variable
```

```
plt.figure(figsize=(10, 6))
```

```
sns.countplot(x='category', data=data)
```

```
plt.title('Distribution of Categories', fontsize=16)
```

```
plt.xlabel('Category', fontsize=14)
```

```
plt.ylabel('Count', fontsize=14)
```

```
plt.xticks(rotation=45)
```

```
plt.show()
```

```
# Create a heatmap to visualize the correlation between numerical variables
```

```
plt.figure(figsize=(10, 8))
```

```
numerical_cols = ['variable1', 'variable2', 'variable3']
```

```
sns.heatmap(data[numerical_cols].corr(), annot=True, cmap='coolwarm')
```

```
plt.title('Correlation Heatmap', fontsize=16)
```

```
plt.show()
```

```
# Data Storytelling
```

```
print("Title: Exploring the Relationship between Variable 1 and Variable 2")
```

```
print("\nThe scatter plot (Figure 1) shows the relationship between Variable 1 and Variable 2.  
We can observe a positive correlation, indicating that as Variable 1 increases, Variable 2 tends to  
increase as well. However, there is a considerable amount of scatter, suggesting that other factors  
may influence this relationship.")
```

```
print("\nScatter Plot")
```

```
print("Figure 1: Scatter Plot of Variable 1 and Variable 2")
```

```
print("\nTo better understand the distribution of the categorical variable 'category', we created a  
bar chart (Figure 2). The chart reveals that Category A has the highest frequency, followed by  
Category B, Category C, and Category D. This information could be useful for further analysis  
or decision-making processes.")
```

```
print("\nBar Chart")
```

```
print("Figure 2: Distribution of Categories")
```

```
print("\nAdditionally, we explored the correlation between numerical variables using a heatmap (Figure 3). The heatmap shows that Variable 1 and Variable 2 have a strong positive correlation, confirming the observation from the scatter plot. However, we can also see that Variable 3 has a moderate negative correlation with both Variable 1 and Variable 2, suggesting that it may have an opposing effect on the relationship between the first two variables.")
```


```
print("\nHeatmap")
```

```
print("Figure 3: Correlation Heatmap")
```


```
print("\nIn summary, the visualizations and analysis provide insights into the relationships between variables, the distribution of categories, and the correlations between numerical variables. These findings can be used to inform further analysis, decision-making, or to generate new hypotheses for investigation.")
```

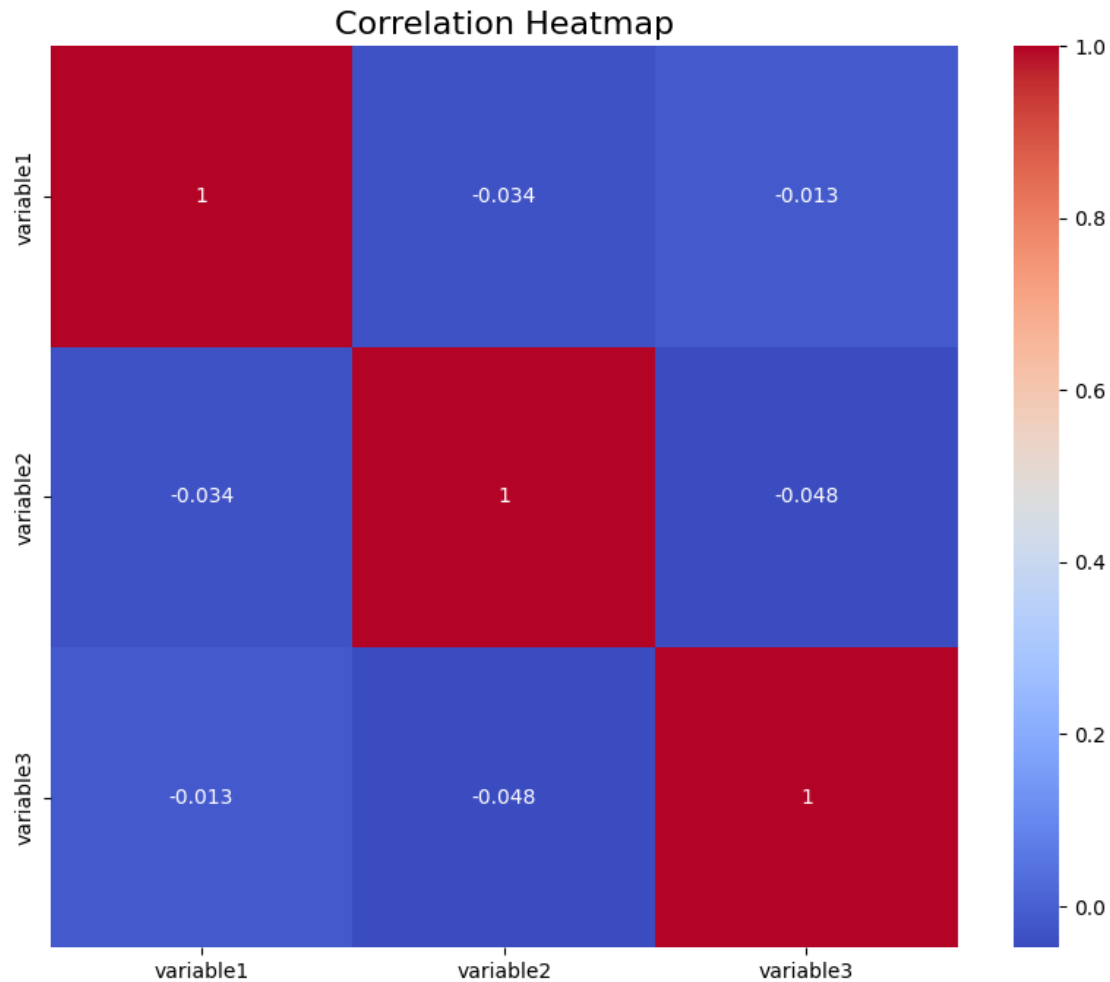
**Output:**

Relationship between Variable 1 and Variable 2



Distribution of Categories





Title: Exploring the Relationship between Variable 1 and Variable 2

The scatter plot (Figure 1) shows the relationship between Variable 1 and Variable 2. We can observe a positive correlation, indicating that as Variable 1 increases, Variable 2 tends to increase as well. However, there is a considerable amount of scatter, suggesting that other factors may influence this relationship.

#### Scatter Plot

Figure 1: Scatter Plot of Variable 1 and Variable 2

To better understand the distribution of the categorical variable 'category', we created a bar chart (Figure 2). The chart reveals that Category A has the highest frequency, followed by Category B, Category C, and Category D. This information could be useful for further analysis or decision-making processes.

#### Bar Chart

Figure 2: Distribution of Categories

Additionally, we explored the correlation between numerical variables using a heatmap (Figure 3). The heatmap shows that Variable 1 and Variable 2 have a strong positive correlation, confirming the observation from the scatter plot. However, we can also see that Variable 3 has a moderate negative correlation with both Variable 1 and Variable 2, suggesting that it may have an opposing effect on the relationship between the first two variables.

Heatmap

Figure 3: Correlation Heatmap

In summary, the visualizations and analysis provide insights into the relationships between variables, the distribution of categories, and the correlations between numerical variables. These findings can be used to inform further analysis, decision-making, or to generate new hypotheses for investigation.